

# Temporally Constrained, Provenance-Verified Evidence Retrieval for Indian Legal Research

Yuvraj Singh

Department of Computer Science and Engineering  
Ajay Kumar Garg Engineering College  
Ghaziabad, India  
yuvraj2312026@akgec.ac.in

Rahul Ranjan

Department of Computer Science and Engineering  
Ajay Kumar Garg Engineering College  
Ghaziabad, India  
rahul2312151@akgec.ac.in

RishiRaj Jaiswal

Department of Computer Science and Engineering  
Ajay Kumar Garg Engineering College  
Ghaziabad, India  
rishi2312003@akgec.ac.in

Saurav Kumar Chaudhary

Department of Computer Science and Engineering  
Ajay Kumar Garg Engineering College  
Ghaziabad, India  
saurav2312023@akgec.ac.in

**Abstract**—AI-assisted legal research fails in ways that fluent output conceals: a cited authority may not exist, may lack the attributed passage, may post-date the matter, or may never have been retrieved. A 2026 Supreme Court of India ruling (2026 INSC 668) overturned tribunal findings built on exactly these defects. We present an evidence pipeline for historical Indian Supreme Court judgments in which every displayed citation is temporally eligible, provenance-bound, and mechanically verifiable, with outcome prediction kept separate. We report three main findings, supported by independent replay and source-side audit. First, naive ILDC-eCourts identifier matching proves unsafe: only 11 of 5,391 syntactic matches survived strict content alignment. Second, pre-ranking temporal eligibility raised Recall@5 from 5/30 to 12/30 and Recall@100 from 12/30 to 15/30 on the frozen 30-case base. Third, all 150 displayed base citations passed verification, with zero unsupported claims, while 15 expected authorities remained absent at  $k=100$ . Outcome baselines on 1,503 cases use their own population.

**Index Terms**—legal information retrieval, explainable AI, provenance verification, temporal filtering, retrieval-augmented systems, Indian legal NLP, BM25, InLegalBERT

## I. Introduction

Legal research rewards fluency over evidence: generated responses cite nonexistent judgments, misattribute passages, or invoke unavailable law. Such hallucinations are frequent, not exceptional [1], and communicating legal standards to models remains open [2].

The problem is no longer confined to the literature. In *Pooja Ramesh Singh v. Jammu and Kashmir Bank Ltd. & Anr.*, 2026 INSC 668 (Civil Appeal No. 11950 of 2025, decided 2 July 2026), the Supreme Court of India set aside tribunal decisions after identifying nonexistent authorities, incorrect citations, and fabricated passages apparently produced through AI-assisted research [3].

We treat that constraint as the design foundation: retrieved facts carry stable locators, duplicates are excluded,

only verbatim passages are rendered, and each citation is verified against corpus and retrieval run.

Our primary finding highlights a stark contrast: while all 150 displayed citations passed strict verification without a single unsupported claim, 15 of the 30 expected authorities failed to appear within the top 100 retrieved results.

**Contributions.** We make four main contributions:

- 1) A quantified cross-corpus identity failure and its remedy. YYYY INSC N to YYYY\_N conversion produced 5,391 syntactic candidates, of which only 11 passed content alignment (Section IV-C).
- 2) Pre-ranking temporal eligibility. Historical availability as a predicate inside the BM25 candidate relation, so ineligible judgments cannot consume retrieval depth (Section VI-D): Recall@5 moved from 5/30 to 12/30 without regressing any prior success.
- 3) A four-state citation protocol with separate denominators. Fabricated, real-but-unretrieved, policy-violating, and valid-but-unexpected citations (Section VII).
- 4) An evaluation that keeps recovery and integrity apart. Authority recovery, citation integrity, and outcome prediction are reported on their own populations (Section VIII).

**Scope.** No legal correctness for retrieved authorities, no autonomous legal advisor, no claim that retrieval improves outcome prediction: a researcher-facing prototype. Two populations are evaluated and never pooled: 1,503 outcome cases and the 30 frozen evidence cases (Base-30).

## II. Related Work

### A. Indian legal NLP and judgment prediction

ILDC established Indian Supreme Court judgment prediction and explanation as a benchmark task [4]; InLe-

galBERT [5] motivates our neural baseline; practitioner-evaluated summarisation [6] shows length and presentation are first-order concerns. None of these lines addresses cross-corpus identity or per-citation provenance, which is where this work sits.

### B. Legal retrieval and temporal constraints

COLIEE, LegalBench, and CaseHOLD evaluate retrieval and reasoning on single fixed corpora with no pre-dating requirement [7]–[9]. TaxFlow adds validity filtering for Indian tax-law QA [10]; CaseFacts and LexTime target evolving validity and event ordering [11], [12]. Retrieval-augmented generation established retrieval as an explicit non-parametric evidence source [13]; BM25 supplies the sparse ranking basis [14] and Tesseract the OCR engine for corpus repair [15].

### C. Positioning

We claim neither the first legal RAG system, nor the first Indian legal model, nor the first temporal legal benchmark: the contribution is the reproducible combination of content-alignment gating, conservative pre-ranking eligibility, exact retrieval-run verification, and separate recovery-versus-validity reporting on a source-first key. Jurisdictions, tasks, corpora, and denominators differ, so no external system is used as a numerical baseline. ILDC [4] supplies splits; InLegalBERT [5] motivates the neural baseline; TaxFlow [10] addresses statutory validity; CaseFacts [11] and LexTime [12] target validity and event ordering.

## III. Problem Formulation

Let a historical query judgment have identifier  $q$ , query year  $Y_q$ , and frozen facts-only text  $F_q$ . Let each candidate evidence passage  $p$  belong to a source judgment with a stable source identifier, an exact decision date, citation metadata, page and character boundaries, and stored passage text.

The temporal policy admits  $p$  only when the decision year of its source is strictly earlier than  $Y_q$ . A same-year source is treated as ambiguous, because ILDC supplies no exact query date; a source with missing or unparseable date metadata is excluded. Both rules trade recall so that, under the fail-closed metadata checks applied here, no displayed authority can post-date the matter.

The evidence task decomposes into three questions. Recovery asks whether the expected authority appears in the top 100 and among the five displayed. Integrity asks whether each citation is a real retrieved passage with exact provenance that avoids the query and duplicates and satisfies the temporal rule. Presentation asks whether the structure aids inspection; we do not evaluate it here.

Outcome prediction is secondary, reported on its own population and never pooled with evidence measures. Table I fixes definitions.

TABLE I  
Operational definitions (frozen evaluation).

| Term                   | Implemented definition                                                                                                                          |
|------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|
| Facts-only input       | Text before the earlier of a recognised dispositive cue and a 60% character cap; inputs below 10% retention or 100 words are excluded           |
| Temporal existence     | The source has a parseable exact decision date                                                                                                  |
| Temporal applicability | For query year $Y$ , an authority is eligible only if its decision year is $< Y$ ; same-year, later-year, and missing-date sources are excluded |
| Temporal effectiveness | Eligibility is applied inside the BM25 candidate relation before ranking and LIMIT 100                                                          |
| Recall@100             | Share of cases whose predefined authority occurs in the returned top-100                                                                        |
| Recall@5               | Share of cases whose predefined authority is among the five selected sources                                                                    |
| Provenance validity    | A displayed item reproduces the stored source ID, citation, decision date, court, locator, exact passage, and retrieval-run membership          |
| Citation groundedness  | Every displayed material proposition is verbatim text of a supplied passage linked to that evidence item                                        |
| Authority consistency  | A displayed source matches the answer-key authority by stable source ID, normalised citation, or title plus exact decision date                 |

### A. Research questions

We investigate two primary research questions. RQ1: Does grounding a legal AI workflow in retrieved evidence improve reliability and retrieval against a facts-only baseline (E2 versus E3; E1 for context)? RQ2: Does provenance-constrained selection and verification reduce unsupported claims?

## IV. Corpus Construction and Auditing

### A. Purpose-split corpus design

We use two corpora for different experimental purposes.

ILDC Single supplies fixed case-level splits and binary outcome labels: 7,593 judgments (5,082 training, 994 validation, 1,517 test). After a shared sufficiency rule excluded 14 test records, the prediction population contains 1,503 cases.

The evidence corpus is an eCourts-derived collection of Indian Supreme Court judgments covering 1950–2020. Of 39,069 English PDF instances, 39,066 accepted PDFs yielded 2,343,435 labelled chunks, of which 2,036,981 unique chunks were loaded into a PostgreSQL provenance store and a SQLite FTS5 BM25 index.

This split follows the information actually available in each source. ILDC supports fixed-split outcome prediction but lacks citation and passage provenance; the eCourts

collection supports dated, traceable retrieval but is not an outcome-label benchmark. Cross-corpus links are drawn only after the alignment checks described below.

#### B. Source-quality audit and targeted OCR repair

A full audit found valid UTF-8 text except 15 PDF instances (14 source IDs) with missing or corrupted embedded text.

Only flagged instances were re-rendered and processed with English Tesseract OCR at 250 DPI [15]. Instances passing a post-OCR quality gate were restored. Raw PDFs were preserved and the index rebuilt from the accepted corpus.

#### C. Cross-corpus identifier alignment

Linking ILDC query cases to eCourts sources requires establishing that two records describe the same judgment. The obvious route — converting an eCourts identifier of the form YYYY INSC N to the ILDC-like string YYYY\_N — is unsafe, and quantifiably so.

That conversion produced 5,391 syntactic candidates. Only 11 passed content alignment; 5,380 were identifier-namespaces collisions.

The corrected pipeline treats identifier equality as candidate generation only. A reusable gate evaluates title/party identity with direct six-token phrase overlap before any mapping drives deduplication or exclusion. Combining syntactic and title/party candidates yielded 8,927 deduplicated pairs (1,304 accepted, 7,623 rejected). Retrieval additionally applies a full-document self-match safeguard, so an unmapped copy of the query judgment is removed without suppressing an earlier authority quoted at length by the later judgment.

#### D. Source-verified authority answer key

The authority evaluation uses 30 ILDC fixed-test cases. For each, the query judgment was inspected through a primary source or accepted eCourts mirror, one earlier authority was recorded independently of system retrieval, and reconciled by stable source ID, normalised citation, or title plus exact date.

The identifier-collision discovery triggered a read-only audit of the then-current mappings: 20 passed, nine failed direct content alignment, and one was unresolved. One case was relinked and the other nine replaced by new fixed-test cases, yielding 30/30 direct-content passes.

The sample is not era-balanced: five query cases are from the 1970s, 13 from the 1980s, nine from the 1990s, two from the 2000s, and one from the 2010s (1980s: 43.3%), a consequence of the source and alignment gates (Section X).

### V. System Design

Four auditable layers implement the design. A corpus and alignment layer links ILDC cases to dated eCourts passages with content-checked deduplication. An outcome branch (E1 sparse linear, E2 InLegalBERT chunk-and-pool) predicts from facts only and never retrieves. An

evidence branch (E3) retrieves BM25 passages under pre-ranking temporal eligibility and renders five source-diverse verbatim passages. A verification layer (E4) checks each displayed record against the corpus and the recorded retrieval run. The renderer is extractive and cannot retrieve, paraphrase material text, or infer outcomes; prediction stays separate.

### VI. Method

#### A. Shared facts-only input

E1 and E2 share facts-only extraction (ilc-predecision-facts-v1): text preceding a dispositive cue and 60% of the document. Cases below 10% retention or 100 words are excluded; selection uses validation only, with one held-out test evaluation.

#### B. E1: sparse linear baseline

E1 uses lowercased TF-IDF unigrams/bigrams (sub-linear TF, min DF two,  $L_2$ , 100,000 features). Of  $C \in \{0.1, 1.0, 10.0\}$ , validation accuracy selected  $C = 10.0$  (smaller  $C$  ties); refit once on eligible training plus validation and evaluated once on test. Seed 202605.

#### C. E2: InLegalBERT chunk-and-pool

E2 fine-tunes law-ai/InLegalBERT (revision b5ecfed8) with a new two-label head over overlapping 512-token windows with 50-token overlap. Training uses three epochs, seed 202607, learning rate  $2 \times 10^{-5}$ , weight decay 0.01, warm-up ratio 0.1, and gradient accumulation.

Mean pooling of window logits before softmax is the primary rule, and checkpoint selection uses validation document-level mean-logit accuracy only (majority vote is a secondary comparison, not a selection rule). All 1,503 eligible test documents and 9,576 test windows are covered.

#### D. E3: temporally constrained retrieval

The query builder (tfidf-segment-salient-terms-v1) emits at most 32 salient terms for a SQLite FTS5 BM25 index backed by PostgreSQL provenance.

The frozen configuration (week11-bm25-salient-terms-preranked-temporal-v3) admits only judgments with decision year strictly earlier than the query year into the candidate relation before BM25 ordering and the top-100 cutoff, so the full depth is spent on eligible material.

Candidate sources are checked against the alignment-gated crosswalk and a self-match rule (100 shared six-token occurrences, 80% source-phrase coverage). A non-learned selector then chooses up to five passages in BM25 order, at most one per source.

#### E. E4: citation and provenance verification

E4 receives the E3 answer, retrieval-run identifier, and persisted corpus records. Each displayed item must exist in the corpus and reproduce its stored passage and provenance. It must also belong to the recorded run, exclude the query and duplicates, and pre-date the query

year; any failure rejects the citation. Authority consistency is then evaluated separately: a fully valid citation may differ from the single expected authority without being substantively irrelevant (Section VIII-D).

#### F. Inference-time evidence augmentation

Outcome prediction was subsequently added to E3 and E4 by reusing the frozen E2 checkpoint without fine-tuning: the frozen facts extract plus the selected verbatim evidence in selection order pass through the same 512-token/50-overlap windows with mean-logit argmax (seed 202607; configuration e3e4-evidence-augmented-inlegalbert-checkpoint6318-v1). Gold labels, scores, dates, and provenance are excluded from model input.

Because the checkpoint was trained on facts-only inputs, applying it to facts plus retrieved passages is an inference-time distribution shift rather than evidence-aware fine-tuning. These predictions are reported descriptively on the frozen 30-case base, and are not directly comparable with the 1,503-case baselines.

#### G. Bundled-intervention caveat

E4 jointly checks corpus existence, passage/provenance identity, run membership, duplicate status, and temporal eligibility. Aggregate reliability cannot be attributed to any one component; the supported claim concerns the complete bundle.

### VII. Citation and Provenance Protocol

Each chunk carries a stable identifier from its source path and passage location, with stored source ID, case ID, citation, exact decision date, court, PDF file, page and character bounds, and chunk text. Retrieval runs record query ID, year, query text, index version, temporal policy, rank, score, and temporal status.

The renderer may expose only supplied selected evidence. Material content remains verbatim; authority cards and conclusions refer to evidence IDs. Verification then checks corpus existence, exact field equality, run membership, duplicate status, and temporal eligibility, with answer-key matching only after a displayed citation has passed. Parallel reporter forms are reconciled through stable source identity or normalised title plus exact decision date, never loose string similarity.

This protocol distinguishes four failure states that a single “correct citation” label would collapse:

- 1) a fabricated or altered citation;
- 2) a real citation that was not retrieved for this query;
- 3) a retrieved citation violating temporal or duplicate policy;
- 4) a fully valid citation that does not match the predefined expected authority.

In our results, states (1)–(3) occur zero times and state (4) occurs in 18 of 30 cases.

TABLE II  
Outcome prediction (distinct populations; not a leaderboard).

| Model                       | $n$   | Acc.     | Macro-F1 |
|-----------------------------|-------|----------|----------|
| Majority class              | 1,503 | 0.5017   | 0.3341   |
| E1 TF-IDF + LR              | 1,503 | 0.61344  | 0.612342 |
| E2 InLegalBERT, mean logits | 1,503 | 0.596806 | 0.592358 |
| E2 InLegalBERT, maj. vote   | 1,503 | 0.6015   | 0.5937   |
| E3 evidence-augmented       | 30    | 0.666667 | 0.603175 |
| E4 verified evid.-augmented | 30    | 0.666667 | 0.603175 |

TABLE III  
RQ1 retrieval and prediction by stratum (Base-30 replayed).

| Population | $n$ | R@5          | R@100        | E3 Acc. | E3 F1  |
|------------|-----|--------------|--------------|---------|--------|
| Base-30    | 30  | 0.40 (12/30) | 0.50 (15/30) | 0.6667  | 0.6032 |

## VIII. Results

### A. Reporting populations

Two populations are reported and never pooled: 1,503 ILDC test cases (outcome prediction) and Base-30 (30 queries; 150 citations) for evidence evaluation.

### B. Outcome prediction

On the 1,503-case population (Table II), sparse E1 reached accuracy 0.61344 and macro-F1 0.612342, exceeding corrected E2 mean-logit pooling at 0.596806 and 0.592358 (majority vote 0.6015 and 0.5937). All primary models exceeded the 0.5017 majority-class accuracy.

While this does not imply that domain pre-training is generally ineffective, it does show that the sparse baseline performed better under our frozen settings.

On the separate 30-case subset, E3 and E4 both reached accuracy 0.666667 and macro-F1 0.603175, with identical confusion matrices.

The equality is expected: with no evidence rejected, the predictor received an identical input.

### C. Authority recovery and evidence integrity

Base-30 recovery is Recall@5 0.40 (12/30) and Recall@100 0.50 (15/30: 12 selected + 3 retrieved-but-unselected); 15 authorities were absent at  $k=100$ .

The principal RQ1 result is the Base-30 row of Table III.

Table IV summarizes Base-30 integrity: 150/150 grounded and provenance-valid with 0/150 temporal violations and 0/30 unsupported claims.

E4 verification is fail-closed: any failing check rejects the citation. On live data nothing was rejected, a ceiling effect of already-valid evidence.

This is the paper’s central result. The system verified everything it displayed (150/150) while recovering the expected authority in 15 of 30 cases at  $k=100$ .

TABLE IV

Evidence recovery and integrity, Base-30 (30 queries; 150 citations).

| Measure             | Value | Arithmetic                      |
|---------------------|-------|---------------------------------|
| Recall@5            | 0.40  | 12 of 30 in top 5               |
| Recall@100          | 0.50  | 15 of 30 in top 100             |
| Absent at $k=100$   | 0.50  | 15 of 30                        |
| Authority precision | 0.08  | 12 of 150 shown                 |
| struct. ceiling     | 0.20  | 30 of 150 (5 shown, 1 credited) |
| Grounding           | 1.00  | 150/150 passed                  |
| Provenance validity | 1.00  | 150/150 passed                  |
| Temporal violations | 0     | 0 of 150                        |
| Unsupported claims  | 0     | 0 of 30 cases                   |

#### D. Interpreting authority-consistency precision

Authority-consistency precision was 0.08, with recall 0.40 and F1 0.133: with one credited authority per case and five displayed, perfect inclusion caps the measure at  $30/150 = 0.20$ , so 0.08 is 40% of ceiling.

The 138 non-matching displayed citations are not independently annotated for substantive legal relevance. They passed every integrity check; they simply are not the one authority the key records.

#### E. Retrieval mechanism

The final configuration came from three recorded development steps on a nine-case probe (0/9 to 3/9 to 6/9 to 7/9 at  $k=100$ ); pre-ranking temporal filtering, carried to the held-out base, moved Recall@5 5/30 to 12/30 and Recall@100 12/30 to 15/30 without losing any of the 12 prior top-100 successes. Probe history is reported from development records, not replayed.

Temporal capacity is not the whole story. 2013\_35 stayed absent despite 79 eligible candidates in its raw top 100, while 1980\_105 was recovered from only three and became a top-five hit: residual lexical mismatch is the prominent unresolved cause.

### IX. Error Analysis

#### A. Outcome disagreement

The analysis is read-only over final frozen outputs. Across 1,503 cases, both models were correct on 684 and both wrong on 368; E1 alone was correct on 238 and E2 alone on 213 (on the 30-case key: 18/3/3/6).

In the evidence-augmented run, E2 was wrong while E3 and E4 were correct for two cases (1974\_36, 1984\_136); E3/E4 verification changed neither the selected evidence nor the shared predictor input (0/30 divergences).

Critically, three cases (2008\_1629, 1981\_187, 1982\_29) retrieved and selected the expected authority yet still produced an incorrect outcome prediction, and a fourth (1985\_40) was retrieved but not selected yet likewise predicted incorrectly: correct recovery does not mechanically imply a correct prediction. No correct prediction was accompanied by an unsupported citation.

#### B. Authority-recovery buckets

The exhaustive partition is: 12 retrieved and selected; 3 retrieved but unselected (1980\_133 rank 15, 1981\_55 rank 28, 1985\_40 rank 78); and 15 absent at  $k=100$ . The middle bucket isolates selection failure (available but undisplayed), while the absent bucket requires upstream improvement.

### X. Limitations and Threats to Validity

Scope. English-language Indian Supreme Court judgments only.

Extraction approximation. Some affected PDF instances were restored through the post-OCR quality gate [15]; repaired text may retain noise. With no gold facts/reasoning boundary, the facts rule is approximate.

Answer-key size and era concentration. Evidence evaluation rests on 30 frozen replayed cases; 13 of the base 30 are from the 1980s. The evaluated set is neither large nor era-balanced and should not be read as representative of the full ILDC test set or of Indian legal research generally.

Single reference authority. One expected authority per query, not every relevant source: inconsistency is a reference mismatch, not a legal-relevance judgment.

Residual recovery gap. At  $k=100$ , 15 of the 30 expected base authorities remain unrecovered; the 2013\_35/1980\_105 contrast shows temporal room alone does not surface the authority.

Year-level granularity. Query dates have year-level precision only, so same-year candidates are conservatively excluded.

Bundled verification. E4 checks integrity properties together and cannot isolate single-component contributions.

Inference numerics. Recomputed E3/E4 mean logits show bounded GPU fp16 variation across runs (bounded at the fourth decimal in the independent recomputation) that changed no decision, evidence selection, or metric; strict serialized file hashes of E3/E4 outputs therefore differ while substantive outputs agree.

Retrieval-equivalence scope. BM25 rebuild equivalence is established at ranking level — identical corpus, parameters, and top-100 agreement on all 30 evaluated queries — not as byte-identical SQLite files, and it does not generalize beyond the evaluated query set.

Prototype scale. Results apply to the frozen implementation and samples. No production deployment is claimed or evaluated.

### XI. Responsible Use and Governance

The system is a legal-research aid, not an adjudicator or provider of legal advice. E3/E4 predictions are experimental; a human reviewer judges relevance, currency, authority, and applicability [3]. Controls are fail-closed and negative results (discarded runs, collision correction, key replacements) are preserved.

## XII. Conclusion and Future Work

We designed, implemented, and evaluated a provenance-preserving, temporally constrained evidence workflow for historical Indian Supreme Court research.

The principal finding is a dissociation. Every one of 150 displayed base citations passed verification, with zero unsupported claims, while 15 of the 30 expected authorities were absent at  $k=100$ .

The cross-corpus audit yields a transferable lesson: only 11 of 5,391 syntactic identifier matches between ILDC and the eCourts collection were content-aligned.

We report supporting results along with their limits. The sparse TF-IDF baseline outperformed the corrected InLegalBERT under the frozen settings (1,503 cases). Inference-time evidence augmentation reached 0.666667 accuracy and 0.603175 macro-F1 on the frozen 30-case subset, descriptively and under acknowledged distribution shift.

Future work. An era-balanced key of 75–100 cases for paired testing; blinded review with independent raters; hybrid dense-sparse retrieval for residual lexical misses; and evidence-sensitive uncertainty language.

### Reproducibility

All configurations, seeds, model revisions, corpus identities, and evaluation artifacts are frozen and machine-readable, including checkpoint-6318 (924a5bb9...dbdc773). The implementation passes 81 tests and the documented checks (14/14 documentation; freeze audit of 39 entries: 28 byte-exact, one line-ending-normalized, 10 metadata-differing). All independently replayed primary metrics matched their frozen references; E3/E4 decisions, evidence selections, and metrics matched despite bounded fp16 variation, and BM25 matched at ranking level on the 30 evaluated queries. The demo stack is covered by 8/8 Spring Boot service tests and 6/6 FastAPI wrapper tests. Code and artifacts are available at <https://github.com/0-YuvrajSingh/NyayaTrace>.

### Acknowledgment

The authors thank Kriti Mishra for her mentorship and guidance throughout this work.

### References

- [1] M. Dahl, V. Magesh, M. Suzgun, and D. E. Ho, “Large Legal Fictions: Profiling Legal Hallucinations in Large Language Models,” *Journal of Legal Analysis*, vol. 16, no. 1, pp. 64–93, 2024.
- [2] J. J. Nay, “Large Language Models as Fiduciaries: A Case Study Toward Robustly Communicating With Artificial Intelligence Through Legal Standards,” *arXiv preprint arXiv:2301.10095*, 2023.
- [3] Supreme Court of India, Pooja Ramesh Singh v. Jammu and Kashmir Bank Ltd. & Anr., 2026 INSC 668 (Civil Appeal No. 11950 of 2025, decided 2 July 2026), 2026. Court PDF: [https://www.sci.gov.in/sci-get-pdf/?diary\\_no=523382025&from=la\\_test\\_judgements\\_order&order\\_date=2026-07-02&type=j](https://www.sci.gov.in/sci-get-pdf/?diary_no=523382025&from=la_test_judgements_order&order_date=2026-07-02&type=j).
- [4] V. Malik, R. Sanjay, S. K. Nigam, K. Ghosh, S. K. Guha, A. Bhattacharya, and A. Modi, “ILDC for CJPE: Indian Legal Documents Corpus for Court Judgment Prediction and Explanation,” in *Proc. 59th Annual Meeting of the Association for Computational Linguistics (ACL-IJCNLP)*, 2021, pp. 4046–4062.
- [5] S. Paul, A. Mandal, P. Goyal, and S. Ghosh, “Pre-trained Language Models for the Legal Domain: A Case Study on Indian Law,” in *Proc. 19th Int. Conf. on Artificial Intelligence and Law (ICAIL)*, 2023, pp. 187–196.
- [6] A. Shukla, P. Bhattacharya, S. Poddar, R. Mukherjee, K. Ghosh, P. Goyal, and S. Ghosh, “Legal Case Document Summarization: Extractive and Abstractive Methods and their Evaluation,” in *Proc. 2nd Conf. of the Asia-Pacific Chapter of the ACL (AACL-IJCNLP)*, 2022, pp. 1048–1064.
- [7] R. Goebel, Y. Kano, M.-Y. Kim, J. Rabelo, K. Satoh, and M. Yoshioka, “Summary of the Competition on Legal Information, Extraction/Entailment (COLIEE) 2023,” in *Proc. 19th Int. Conf. on Artificial Intelligence and Law (ICAIL)*, 2023, pp. 472–480.
- [8] N. Guha et al., “LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models,” in *Advances in Neural Information Processing Systems 36, Datasets and Benchmarks Track*, 2023.
- [9] L. Zheng, N. Guha, B. R. Anderson, P. Henderson, and D. E. Ho, “When Does Pretraining Help? Assessing Self-Supervised Learning for Law and the CaseHOLD Dataset of 53,000+ Legal Holdings,” in *Proc. 18th Int. Conf. on Artificial Intelligence and Law (ICAIL)*, 2021, pp. 159–168.
- [10] V. R. Karna, R. R. M., B. S. Babu, et al., “A Hybrid RAG-LLaMA Framework for Scalable and Accurate Interpretation of Legal Texts,” *Applied Artificial Intelligence*, vol. 40, no. 1, art. 2626097, 2026. DOI: <https://doi.org/10.1080/08839514.2026.2626097>.
- [11] A. R. Putta, J. Devasier, and C. Li, “CaseFacts: A Benchmark for Legal Fact-Checking and Precedent Retrieval,” in *Proc. 64th Annual Meeting of the Association for Computational Linguistics*, 2026, pp. 17246–17265.
- [12] C. Barale, L. Barrett, V. S. Bajaj, and M. Rovatsos, “LexTime: A Benchmark for Temporal Ordering of Legal Events,” in *Findings of the Association for Computational Linguistics: EMNLP 2025*, 2025, pp. 5220–5236.
- [13] P. Lewis et al., “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks,” in *Advances in Neural Information Processing Systems 33*, 2020, pp. 9459–9474.
- [14] S. Robertson and H. Zaragoza, “The Probabilistic Relevance Framework: BM25 and Beyond,” *Foundations and Trends in Information Retrieval*, vol. 3, no. 4, pp. 333–389, 2009.
- [15] R. Smith, “An Overview of the Tesseract OCR Engine,” in *Proc. 9th Int. Conf. on Document Analysis and Recognition (ICDAR)*, vol. 2, 2007, pp. 629–633.